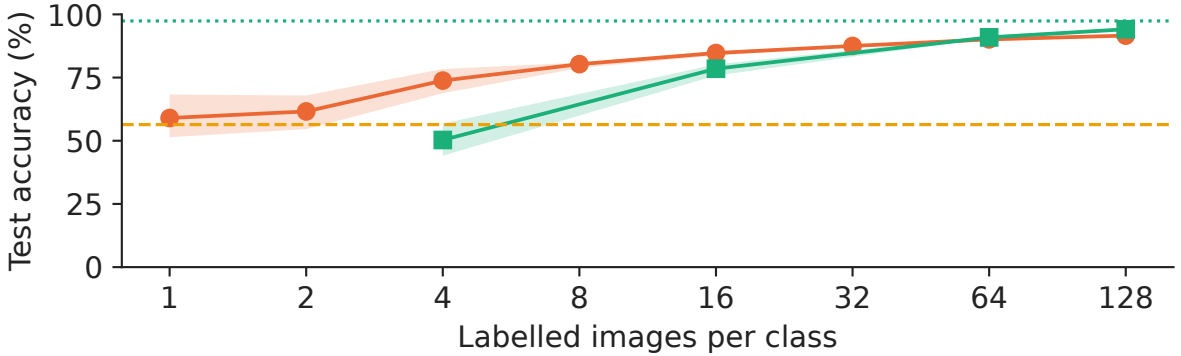
Test accuracy (%)
100
75
50
25
0
1
2
4
8
16
32
64
128
Labelled images per class
CLIP ViT-B/16 embeddings + logistic (5 seeds)
CNN trained from scratch (3 seeds)
CLIP zero-shot (no labelled image)
CNN on all 21,600 training images